

Clojure, Concurrency, And You

Abhinav Sarkar

(((((IN/Clojure Bengaluru 2019))))))

Time

```
(defn pass-time [person]
  (-> person
    (assoc :hair-color :gray)
    (update :age inc)))
```

Values and Identities

```
user=> (def abhinav {:hair-color :black :age 33})  
#'user/abhinav
```

Values and Identities

```
user=> (def abhinav {:hair-color :black :age 33})  
#'user/abhinav
```

```
user=> (def abhinav  
  #_=> (atom {:hair-color :black :age 33}))  
#'user/abhinav  
user=> (swap! abhinav pass-time)  
{:hair-color :gray, :age 34}
```

Concurrency

Concurrency is a program-structuring technique in which there are multiple threads of control which execute "at the same time".

— Simon Marlow, Parallel and Concurrent Programming in Haskell

Threads

- » Thread is a sequence of instructions along with a context.
- » In case of Clojure on JVM, the threading model is provided by the JVM which only supports OS threads.


```

fs.readdir(source, function (err, files) {
  if (err) {
    console.log('Error finding files: ' + err);
  } else {
    files.forEach(function (filename, fileIndex) {
      console.log(filename);
      gm(source + filename).size(function (err, values) {
        if (err) {
          console.log('Error identifying file size: ' + err);
        } else {
          console.log(filename + ' : ' + values);
          aspect = (values.width / values.height);
          widths.forEach(function (width, widthIndex) {
            height = Math.round(width / aspect);
            console.log('resizing ' + filename + 'to ' + height + 'x' + height);
            this.resize(width, height).write(dest + 'w' + width + '_' + filename,
              function(err) {
                if (err)
                  console.log('Error writing file: ' + err);
              });
          }.bind(this));
        }
      });
    });
  }
});

```

```
user=> (def t (Thread. #(println "hello")))
#'user/t
user=> (.start t)
hello
nil
```

```
user=> (import java.util.concurrent.Executors)
java.util.concurrent.Executors
user=> (def tp (Executors/newSingleThreadExecutor))
#'user/tp
user=> (.submit tp #(println "hello"))
hello
#object[java.util.concurrent.FutureTask ...]
```

```
user=> (future (println "hello"))  
hello  
#object[clojure.core$future_call$reify__6962 ...]
```

```
user=> (future (println "hello"))
hello
#object[clojure.core$future_call$reify__6962 ...]
```

```
user=> (def f (future (do
  #_=> (println "hello")
  #_=> 12345)))
hello
#'user/f
user=> (deref f)
12345
```

Synchronization

Synchronization

- » The process by which multiple threads agree on *some things* at *some time*.
- » For example:
 - » timing: forking and joining threads
 - » value of a variable
 - » a sequence of steps to execute
 - » access to a shared resource

√ Time / Things >	One value	Multiple values
Synchronous	Lock, Atom	Multiple locks, Ref
Asynchronous	Agent	CRDTs, Raft/Paxos

Locks

Locks

- » An easy way of synchronization.
- » Prevent concurrent access to critical sections/memory.
- » Do not compose.

```
user=> (def lock (Object.))  
#'user/lock  
user=> (locking lock  
      #_=> (println "locked hello"))  
locked hello  
nil
```

```
user=> (import java.util.concurrent.locks.ReentrantLock)
java.util.concurrent.locks.ReentrantLock
user=> (def lock (ReentrantLock.))
#'user/lock
user=> (try
  #_=>    (.lock lock)
  #_=>    (println "locked hello")
  #_=>    (finally
  #_=>      (.unlock lock)))
locked hello
nil
```

Atoms

Atoms

- » Atoms are references which change atomically and immediately.
- » Simplest of all reference types.
- » Do not compose.

```
user=> (def abhinav
  #_=> (atom {:hair-color :black :age 33}))
#'user/abhinav
user=> (swap! abhinav pass-time)
{:hair-color :gray, :age 34}
user=> (reset! abhinav {:hair-color :none :age 33})
{:hair-color :none, :age 33}
user=> @abhinav
{:hair-color :none, :age 33}
```

```
package clojure.lang;
import java.util.concurrent.atomic.AtomicReference;

final public class Atom {
    private final AtomicReference state;

    public Atom(Object o) { state = new AtomicReference(o); }

    public Object deref() { return state.get(); }

    public Object swap(IFn f) {
        for (;;) {
            Object v = deref();
            Object newv = f.invoke(v);
            if (state.compareAndSet(v, newv)) {
                notifyWatches(v, newv);
                return newv;
            }
        }
    }
}
```


Atom

- » Most ubiquitous concurrency feature used in Clojure.
- » Use cases: dynamic configs, database connections, simple caches.
- » Do not call swap with long running or non-idempotent functions.

Agents

Agents

- » Agents, like Atoms, are references which support atomic changes.
- » But the changes are made in an asynchronous fashion.
- » Do not compose.

```
user=> (def counter (agent 0))
#'user/counter
user=> (dotimes [i 10]
      #_=> (send counter inc))
nil
user=> (await counter)
nil
user=> (println @counter)
10
nil
```

Agents are not Actors

```
-module(counter).  
-export([loop/1]).
```

```
loop(N) -> receive  
  {inc} -> loop(N+1);  
  {get, Sender} -> Sender ! N, loop(N)  
end.
```

```
> Pid = spawn(counter, loop, [0]).  
> Pid ! {inc}.  
> Pid ! {get, self()}.  
> receive Value -> io:fwrite("~p~n", [Value]) end.  
1
```

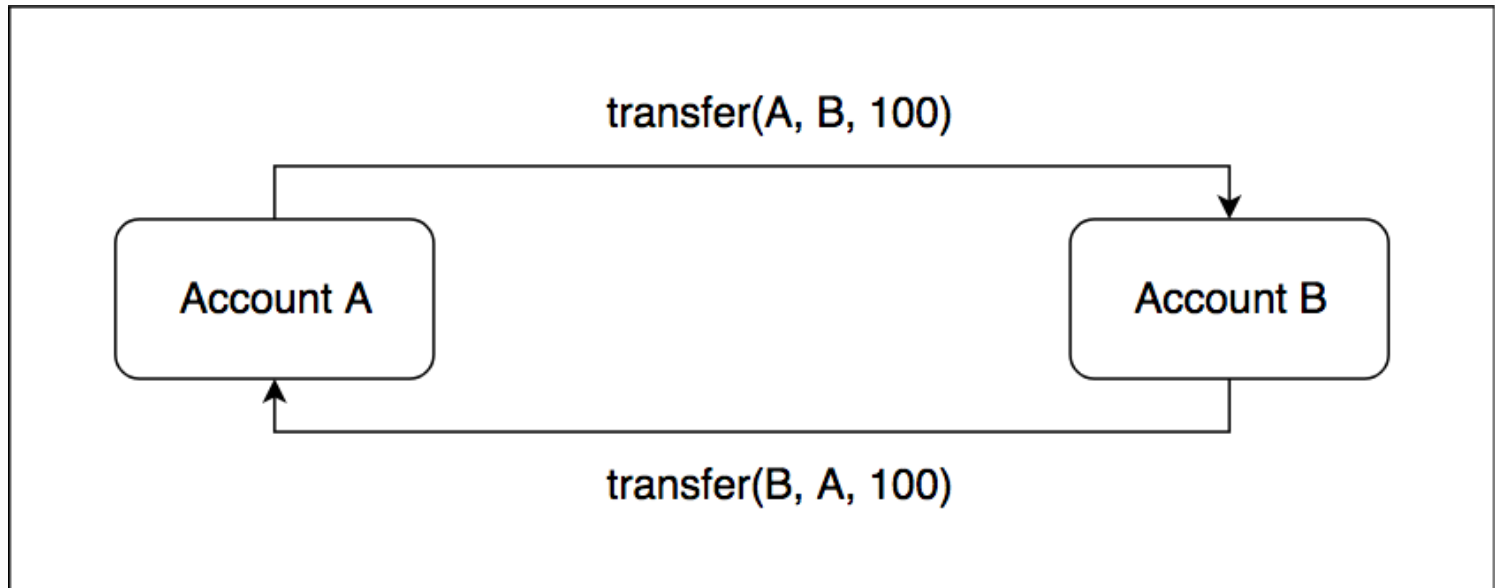
Agents

- » Can be used for any state that does not require strict consistency for reads:
 - » Counters (e.g. message rates in event processing)
 - » Collections (e.g. recently processed events)
- » Can be used for offloading arbitrary computations to a thread pool using `send-via`.
- » Uses an unbounded queue, so too many functions enqueued in it may cause OOM.

Refs

Refs and Software Transactional Memory

- » Refs allows changing multiple references together in a single atomic operation.
- » **Atomicity:** All the state changes become visible to all the threads at once.
- » **Consistency:** All the state changes can be validated before allowing the transaction to commit.
- » **Isolation:** The atomic operation is completely unaffected by whatever other threads are doing.



```
import java.util.concurrent.locks.Lock;
import java.util.concurrent.locks.ReentrantLock;

class Account {
    private int id, amount;
    private Lock lock = new ReentrantLock();

    Account(int id, int initialAmount) {
        this.id = id;
        this.amount = initialAmount;
    }

    public void withdraw(int n) {
        this.lock.lock();
        try { this.amount -= n; }
        finally { this.lock.unlock(); }
    }

    public void deposit(int n) { this.withdraw(-n); }

    public void transfer(Account other, int n) {
        this.withdraw(n);
        other.deposit(n);
    }
}
```

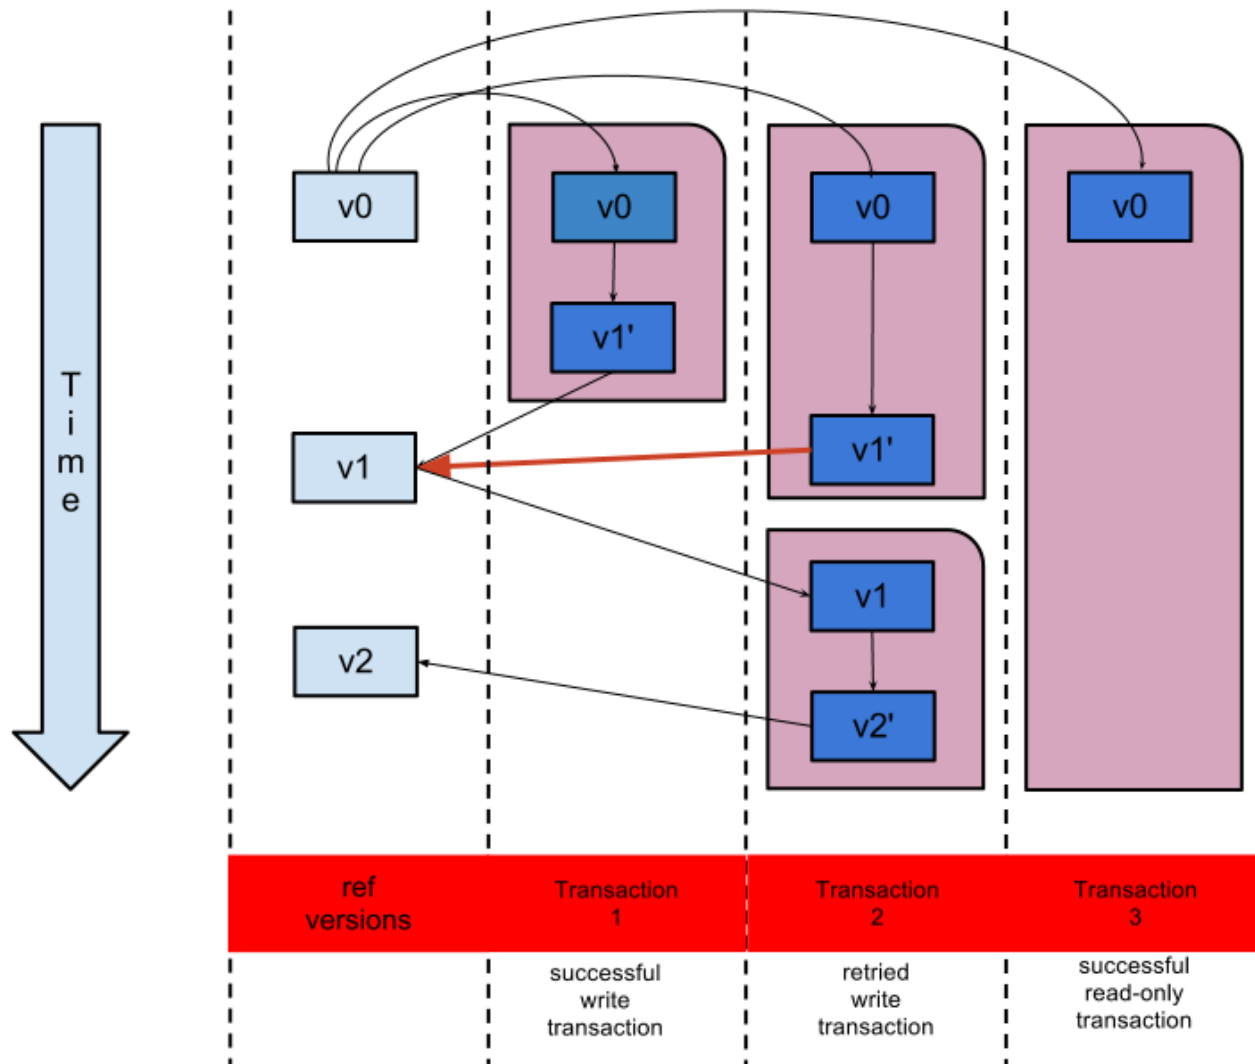
```
public void transfer(Account other, int n) {  
    if (this.id < other.id) {  
        this.lock.lock();  
        other.lock.lock();  
    } else {  
        other.lock.lock();  
        this.lock.lock();  
    }  
    try {  
        this.amount -= n;  
        other.amount += n;  
    } finally {  
        if (this.id < other.id) {  
            this.lock.unlock();  
            other.lock.unlock();  
        } else {  
            other.lock.unlock();  
            this.lock.unlock();  
        }  
    }  
}
```

```
(def account1 (ref 100))
(def account2 (ref 100))


(defn withdraw [account amount]
  (alter account - amount))

(defn deposit [account amount]
  (withdraw account (- amount)))

(defn transfer [from to amount]
  (dosync
    (withdraw from amount)
    (deposit to amount)))
```



t: 1 4 7



Clojure STM vs. Haskell STM

```
import System.IO
import Control.Concurrent.STM

type Account = TVar Int

withdraw :: Account -> Int -> STM ()
withdraw acc amount = do
    bal <- readTVar acc
    writeTVar acc (bal - amount)

deposit :: Account -> Int -> STM ()
deposit acc amount = withdraw acc (- amount)

transfer :: Account -> Account -> Int -> IO ()
transfer from to amount = atomically $ do
    deposit to amount
    withdraw from amount
```

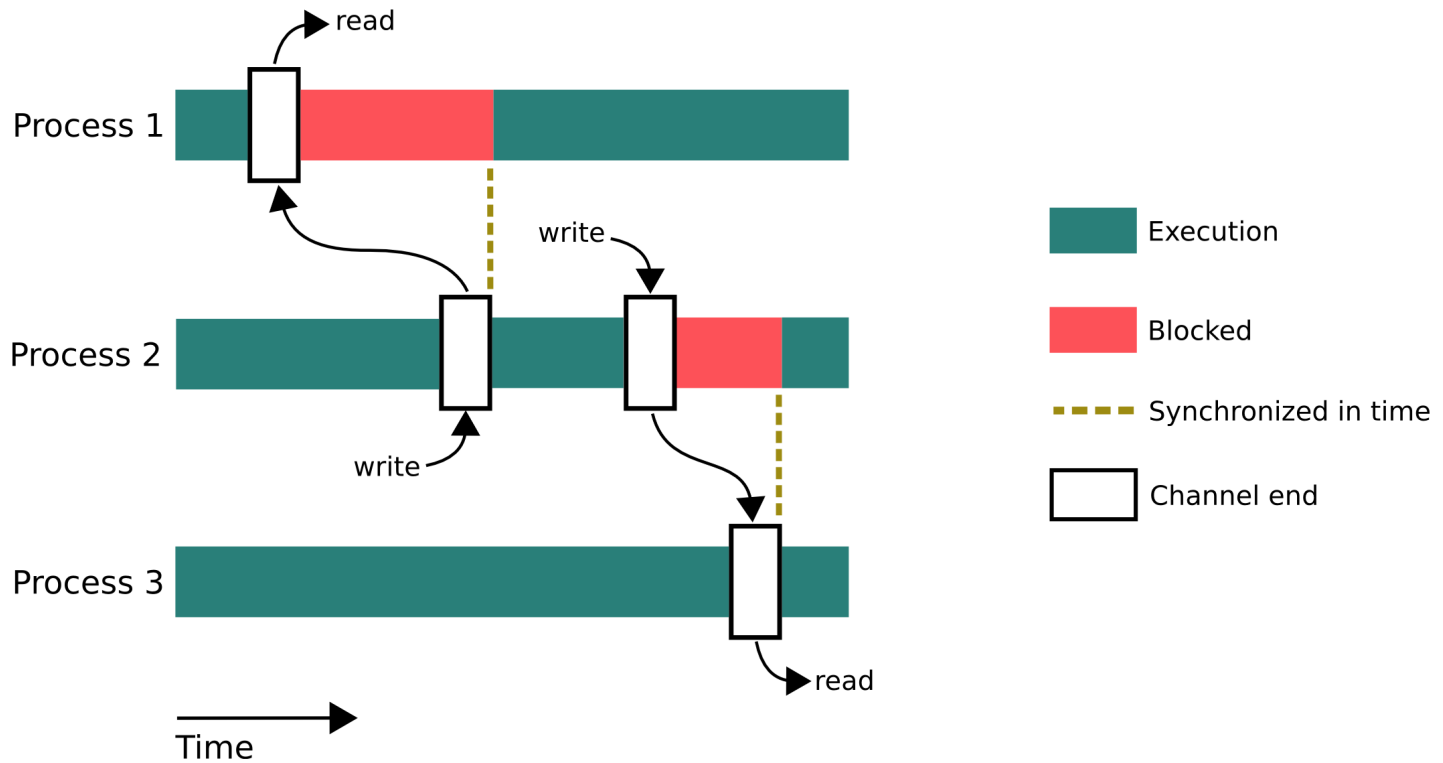
Refs

- » Best option for implementing in-memory transaction data stores
 - » Chat servers
 - » Multiplayer games
- » In-memory stream computation solutions.
- » A long-running transaction may re-execute many times because it may be repeatedly aborted by shorter transactions.
- » Keeping history of values is expensive. Even more so when the values are not persistent collections.

core.async

core.async and Communicating Sequential Processes

- » Independent threads of activity.
- » Synchronous communication through channels.
- » Multiplexing of channels with alternation.



```
(go (println "hi"))
```

```
(def echo-chan (chan))
```

```
(go (println (<! echo-chan)))
```

```
(go (>! echo-chan "hello"))
```

```
; => hello
```

```
(def echo-chan (chan 10))
```

```
(let [c1 (chan)
      c2 (chan)
      c3 (chan)]
  (dotimes [n 3]
    (go
      (let [[v ch] (alts! [c1 c2 c3])]
        (println "Read" v "from" ch))))
  (go (>! c1 "hello")))
(go (>! c2 "allo"))
(go (>! c2 "hola")))
; => Read allo from #<ManyToManyChannel ...>
; => Read hola from #<ManyToManyChannel ...>
; => Read hello from #<ManyToManyChannel ...>
```

Multiplexing

```
user=> (def input (chan 1))
#'user/input
user=> (def broadcast (mult input))
#'user/broadcast
user=> (dotimes [i 3]
  #_=> (let [output (chan 1)]
    #_=> (tap broadcast output)
    #_=> (go-loop []
      #_=> (if-let [v (<! output)]
        #_=> (do (println i "Got!" v)
          (recur))
        (println "Exiting")))))
nil
```

Multiplexing

```
user=> (>!! input 42)
```

```
true0 Got! 421
```

```
2 Got! Got!42
```

```
42
```

```
user=> (>!! input 43)
```

```
true
```

```
01 Got!Got! 4343
```

```
2
```

```
Got! 43
```

```
user=> (close! input)
```

```
Exitingnil
```

```
Exiting
```

```
Exiting
```

Publish-Subscribe

```
user=> (def input (chan 1))
#'user/input
user=> (def p (pub input :tag))
#'user/p
user=> (let [c (chan 1)]
  #_=> (sub p :cats c)
  #_=> (go-loop []
    #_=> (if-let [v (<! c)]
      #_=> (println "Cat guy got:" v)
      #_=> (println "Cat guy exiting"))))
#object[clojure.core.async.impl.channels.ManyToManyChannel ...]
user=> (let [c (chan 1)]
  #_=> (sub p :dogs c)
  #_=> (go-loop []
    #_=> (if-let [v (<! c)]
      #_=> (println "Dog guy got:" v)
      #_=> (println "Dog guy exiting"))))
#object[clojure.core.async.impl.channels.ManyToManyChannel ...]
```


Publish-Subscribe

```
user=> (defn send-with-tags [msg]
  #_=> (doseq [tag (:tags msg)]
    #_=> (println "sending..." tag)
    #_=> (>!! input {:tag tag :msg (:msg msg)})))
#'user/send-with-tags
user=> (send-with-tags {:msg "New Cat Story" :tags [:cats]})
sending... :cats
nil
Cat guy got: {:tag :cats, :msg New Cat Story}
user=> (send-with-tags {:msg "New Dog Story" :tags [:dogs]})
sending... :dogs
nil
Dog guy got: {:tag :dogs, :msg New Dog Story}
```

Publish-Subscribe

```
user=> (let [c (chan 1)]
  #_=> (sub p :dogs c)
  #_=> (sub p :cats c)
  #_=> (go-loop []
  #_=> (if-let [v (<! c)]
  #_=> (println "Cat/Dog guy got:" v)
  #_=> (println "Cat/Dog guy exiting"))))
#object[clojure.core.async.impl.channels.ManyToManyChannel]
```

Pipelines

```
(pipeline n to xf from close? ex-handler)
```

```
(pipeline-async n to af from)
```

```
(pipeline-blocking n to xf from close? ex-handler)
```

core.async vs Goroutines

```
func fibonacci(c, q chan int) {  
    x, y := 0, 1  
    for {  
        select {  
        case c <- x:  
            x, y = y, x+y  
        case <-q:  
            fmt.Println("quit")  
            return  
        }  
    }  
}
```

```
func main() {  
    c := make(chan int)  
    quit := make(chan int)  
    go func() {  
        for i := 0; i < 10; i++ {  
            fmt.Println(<-c)  
        }  
        quit <- 0  
    }()  
    fibonacci(c, quit)  
}
```

core.async

- » Data transformation pipelines like ETL.
- » Multi-user chat servers, game servers.
- » More broadly, Staged Event Driven Architecture programs.
- » With Clojurescript, it is a great replacement for callback for UI interactions.

core.async

- » Data transformation pipelines like ETL.
- » Multi-user chat servers, game servers.
- » More broadly, Staged Event Driven Architecture programs.
- » With Clojurescript, it is a great replacement for callback for UI interactions.
- » Doing blocking IO in go-threads blocks them.
- » Error handling is complicated.
- » Too many pending puts or take may throw errors.



References

- » [Clojure reference documentation](#)
- » [Clojure STM: What, why, how?](#)
- » [Thoughts on STM](#)
- » [Implementation details of core.async Channels](#)
- » [Core Async Go Macro Internals](#)
- » [Publish and Subscribe with core.async](#)

Thanks

[@abhin4v](#)

abnv.me/ccy

